

Structured-Grid SOR with Temporal Blocking: from one Xeon thread to a CUDA-tiled L40S

Che-Yu Yu
EC527, Spring 2026

April 29, 2026

1 Description of the Algorithm

1.1 Mathematical foundation

We are solving the steady-state, source-free 2D Laplace equation on the unit square with Dirichlet boundary data:

$$\nabla^2 u(x, y) = \frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} = 0, \quad u|_{\partial\Omega} = g(x, y). \quad (1)$$

Discretise on a uniform $N \times N$ grid with spacing $h = 1/(N-1)$. The standard second-order centred difference

$$\frac{u_{i-1,j} - 2u_{i,j} + u_{i+1,j}}{h^2} + \frac{u_{i,j-1} - 2u_{i,j} + u_{i,j+1}}{h^2} = 0$$

rearranges to the canonical 5-point stencil shown in Figure 1:

$$u_{i,j} = \frac{1}{4}(u_{i-1,j} + u_{i+1,j} + u_{i,j-1} + u_{i,j+1}). \quad (2)$$




Figure 1: The 5-point Laplacian stencil. Each interior cell averages its four orthogonal neighbours.

Equation (2) is the *Jacobi update*: cell (i, j) 's next value depends only on its neighbours' *current* values. Treating it as a fixed-point iteration $u^{(k+1)} = \frac{1}{4}(\dots)$ and substituting into (1) as the residual, Successive Over-Relaxation (SOR) adds an over-relaxation parameter ω to the Gauss-Seidel update:

$$u_{i,j}^{(k+1)} = (1 - \omega) u_{i,j}^{(k)} + \frac{\omega}{4} \left(u_{i-1,j}^{(k+1)} + u_{i+1,j}^{(k)} + u_{i,j-1}^{(k+1)} + u_{i,j+1}^{(k)} \right). \quad (3)$$

Because the right-hand side mixes already-updated neighbours $(u_{i-1,j}^{(k+1)}, u_{i,j+1}^{(k+1)})$ with not-yet-updated ones $(u_{i+1,j}^{(k)}, u_{i,j-1}^{(k)})$, each cell is using the freshest information available at the moment it's visited. That's the whole trick: Gauss-Seidel converges roughly $2 \times$ faster than Jacobi, and SOR with the optimal ω converges $O(N)$ times faster than Jacobi for the 2D Laplacian.

Why SOR converges so much faster. Plugging the iteration matrix into the spectral-radius analysis (Briggs–Henson–McCormick, *A Multigrid Tutorial*, ch. 2) gives, for the Poisson problem on a uniform grid,

$$\omega_{\text{opt}} = \frac{2}{1 + \sin(\pi/(N-1))} \approx 2 - \frac{2\pi}{N}, \quad (4)$$

and the spectral radius (a.k.a. asymptotic convergence factor) drops from $\rho_{\text{Jacobi}} = \cos(\pi/(N-1)) \approx 1 - \pi^2/(2N^2)$ to $\rho_{\text{SOR}} = 1 - \frac{2\pi}{N}$. In iteration-count terms, that’s $O(N^2)$ iterations for Jacobi vs $O(N)$ for tuned SOR. The single biggest algorithmic optimisation in this whole project, and we verify it experimentally in 2.5.2 (E08). At $N=128$ the predicted ratio is $\sim 13\times$; we measured $13.5\times$. Hopefully the math will be that nice every time.

1.2 The specific algorithmic form used

We implemented three forms of the sweep. The *baseline* is a straight scan-order rewrite of (3):

```
/* baseline SOR, scan order, in-place: this is sequential */
for (i = 1; i < N-1; i++)
  for (j = 1; j < N-1; j++) {
    sum = u[i-1][j] + u[i+1][j] + u[i][j-1] + u[i][j+1];
    u[i][j] = (1.0f - OMEGA) * u[i][j] + (OMEGA * 0.25f) * sum;
  }
```

The trouble: when we write to $u[i][j]$, the next column update at $u[i][j+1]$ reads $u[i][j]$. The just-overwritten value. And the next row reads $u[i+1][j]$, which was overwritten one row ago. This is the classical *loop-carried dependency* that makes the scan-order SOR sweep strictly sequential. We can’t slap an `#pragma omp parallel for` on this and call it a day; it would race.

Red-black SOR: the parallel-friendly cousin. Colour the grid like a checkerboard. “Red” cells $(i+j)$ even depend only on “black” cells $(i+j)$ odd, and vice versa. So an entire colour can be updated independently in parallel:

```
/* red-black SOR: two colour-passes; each pass is embarrassingly parallel */
for (parity = 0; parity < 2; parity++) {
  #pragma omp parallel for
  for (i = 1; i < N-1; i++) {
    int j0 = ((i + parity) & 1) ? 1 : 2;
    for (j = j0; j < N-1; j += 2) {
      sum = u[i-1][j] + u[i+1][j] + u[i][j-1] + u[i][j+1];
      u[i][j] = (1.0f - OMEGA) * u[i][j] + (OMEGA * 0.25f) * sum;
    }
  }
}
```

Two half-sweeps per “iteration”; each half-sweep is a parallel-for. We use this in `sor2d_rb.c` for the omega-tuning study (2.5.2, E08).

Jacobi ping-pong: the variant we actually parallelise. Most of the project uses the still-simpler form: read from `src`, write to `dst`, swap. No within-sweep dependency at all. Every cell is computed from the previous sweep’s data:

```
/* Jacobi ping-pong: dst[i][j] = src[i][j] - OMEGA * (src[i][j] - avg) */
#pragma omp parallel for
for (i = 1; i < N-1; i++)
  for (j = 1; j < N-1; j++) {
```

```

    nb = 0.25f * (src[i-1][j] + src[i+1][j] + src[i][j-1] + src[i][j+1]);
    dst[i][j] = src[i][j] - OMEGA * (src[i][j] - nb);
}
swap(&src, &dst);

```

We trade convergence rate (this form is unconditionally stable only for $\omega \in (0, 1]$ on large grids. We use $\omega=0.9$) for clean parallelism and the ability to do *temporal blocking* cleanly.

Temporal blocking (the shrinking trapezoid). For each *super-step* of T Jacobi sweeps:

1. Tile the grid into $B \times B$ output regions (Figure 2).
2. For each tile, load a $(B+2T)^2$ scratch buffer from `src` with clamp-to-edge padding.
3. Run T in-scratch Jacobi sweeps; the valid update region shrinks by one cell on each side per sub-step (Figure 2). After T sub-steps the central B^2 holds the correct T -step-advanced state.
4. Write the central B^2 back to `dst`.

DRAM traffic for the central cells drops by a factor of T , since we read them once per super-step instead of once per sweep. The cost is redundant halo work: the scratch volume is $(B+2T)^2$, the output volume is B^2 , so the work multiplier is $((B+2T)/B)^2$. That's the central tradeoff and we'll measure its unimodal optimum in 2.5 (E04).




Figure 2: Spatial vs temporal blocking. Same grid, two decompositions.

Variable glossary used throughout:

- N . Grid side (interior is $(N-2)^2$ cells; edges held at Dirichlet boundary).
- B . Spatial block (tile) size. L1/L2 budget determines the upper bound; 64-cell minimum to amortise tile overheads.
- T . Temporal halo / super-step depth. Number of sub-sweeps fused per DRAM round-trip.
- ω . Relaxation factor. 0.9 for Jacobi-ping-pong kernels, $\omega_{\text{opt}} \approx 2 - 2\pi/N$ for the red-black SOR variant.
- `iters`. Total sweep count. Must satisfy `iters % T == 0` for the temporal kernels.
- `TOL`. Convergence tolerance (10^{-5} in 2.5.2).

1.3 Implementation reality check

Before writing a single optimization, it’s worth working through a small example by hand to make sure the math is what we think it is, and to surface the FP precision question early.

A worked numerical example. Take a 5×5 grid with fixed Dirichlet boundary $u \equiv 1$ and interior init $u \equiv 0$. The true continuous solution to Laplace’s equation with constant boundary $u \equiv 1$ is $u \equiv 1$ everywhere (the discrete one is the same to machine precision; the iteration is just there to *find* it from a wrong initial guess). Run scan-order SOR with $\omega = 1.6$ on the 3×3 interior. After sweep 1:

$$u^{(1)} = \begin{bmatrix} \mathbf{1} & 1 & 1 & 1 & \mathbf{1} \\ \mathbf{1} & 0.8000 & 0.7200 & 1.0880 & \mathbf{1} \\ \mathbf{1} & 0.7200 & 0.5760 & 1.0656 & \mathbf{1} \\ \mathbf{1} & 1.0880 & 1.0656 & 1.6525 & \mathbf{1} \\ \mathbf{1} & 1 & 1 & 1 & \mathbf{1} \end{bmatrix}$$

(boundary in bold; interior in normal weight). The (3,3) corner overshoots to 1.6525. Way past the true value of 1. And that’s exactly what over-relaxation looks like: we’re betting that overshooting now will save us iterations later. After sweep 2 the center cell is at 1.3005, sweep 3 has it ringing back below 1, and by about sweep 30–40 the whole grid settles to within 10^{-5} of the true $u \equiv 1$. We’ll see this U-curve again in 2.5.2, where we sweep ω explicitly and find ω_{opt} matches (4) to within 0.1%.

The honest tension. Equations (3) and the loop above are the textbook *sequential* SOR. Once we want parallelism (which we do, this is a parallel-programming course), the loop-carried dependency along the scan order forbids any naive `#pragma omp parallel for` on the outer `i` loop: two threads racing on adjacent rows would have one thread reading `u[i-1][j]` before the other thread has written to it. We have two ways out:

1. **Red-black SOR.** Colour the grid; update each colour in parallel; do two half-sweeps per “iteration”. ω_{opt} shifts slightly because the iteration matrix is different, but the asymptotic $O(N)$ behaviour is preserved. We use this in `sor2d_rb.c` for the omega-tuning study.
2. **Damped Jacobi ping-pong.** Read `src`, write `dst`, swap. Pure parallelism inside one sweep but $\omega \in (0, 1]$ stability constraint, so we lose the over-relaxation factor and pay $O(N^2)$ iterations to converge. We use this in all the temporal-blocking kernels because the shrinking trapezoid in Figure 2 works cleanly with a read-only `src`.

So we won’t be running “classical” SOR through the GPU. We’ll run damped-Jacobi ping-pong with temporal blocking through the GPU, and quantify the iter-count loss separately by sweeping ω in the red-black variant in 2.5.2. The honest apples-to-apples comparison at the end is per-iter throughput (CPE), not iters-to-tolerance.

Validation strategy. Three layers, used everywhere:

1. Every CPU temporal kernel is checked *bit-identical* against its in-binary baseline (Jacobi ping-pong is deterministic; reordering tiles must not change the result modulo FMA contraction). Every CPU smoke run prints `max|base-temp| = 0.0000e+00`.
2. The GPU temporal kernel validates against the GPU baseline at $\sim 1.4 \times 10^{-7}$ relative. Pure FP32 round-off accumulating over 96 sweeps; the bound is $\sqrt{N \cdot \text{iters}} \cdot \varepsilon_{\text{FP32}} \approx 10^{-5}$, comfortably under our TOL. Both GPU kernels disagree with the FP64 CPU reference at the same magnitude.

3. The omega-sweep in 2.5.2 reproduces $\omega_{\text{opt}} = 1.9517$ to 0.1% on $N=128$, which is by itself a sanity check that the discretisation, boundaries, and convergence detector are correctly wired up.

2 Scalar CPU Version

2.1 The hot inner loop

The 5-point stencil is the only thing worth optimising. It's executed $N^2 \cdot \text{iters}$ times. For $N=2048$ and 64 iters that's a billion times. And the rest of the program (boundary copies, allocator calls, the main loop) is sublinear at N^2 or $O(\text{iters})$.

Per cell-update: we have 5–7 floating-point ops (4 adds for the neighbour sum, 1 multiply for $\frac{1}{4}$, plus the ω -update which is either an FMA or a multiply-and-subtract depending on the form), against 1 read of `src[i][j]` and 1 write to `dst[i][j]` at *DRAM granularity*. The four neighbours stream through the same cache lines as we move along the row, so they don't cost any extra DRAM traffic in steady state. At FP32, that's 8 B/cell of DRAM traffic and roughly 6 FLOPs/cell, giving an arithmetic intensity of

$$\text{AI} \approx \frac{6 \text{ FLOPs}}{8 \text{ bytes}} = 0.75 \text{ FLOP/byte.} \quad (5)$$

That's deep in the bandwidth-bound regime (Figure 3). The roofline ridge for our Xeon Gold 6426Y socket sits near ~ 5 FLOP/byte; the L40S sits near ~ 100 FLOP/byte. At 0.75 we're firmly on the diagonal. *Every* optimization in this project is therefore a *memory* optimization in disguise: we are hunting for ways to reduce DRAM bytes per useful FLOP. The arithmetic itself is already free.




Figure 3: Roofline-style throughput curve for our CPU configuration. The 5-point stencil's $\text{AI} \approx 0.5\text{--}0.75$ puts us where the red arrow lands. On the bandwidth-bound diagonal, well below the peak- compute plateau. This is real measured data, not the canonical double-roof cartoon: the curve is what we observe as we vary the working-set size and AI together.

Two micro-variants of the inner loop. We tried both *ij*-order (the textbook scan-order) and *ji*-order (reversed) on the inner loop, with and without `restrict` qualifiers on `src` and `dst`. The numbers at $N=512$, single-thread, 64 iters:

variant	CPE
ij-order, no restrict	0.396
ij-order, with <code>restrict</code>	0.392 (-1%)
ji-order, with <code>restrict</code>	1.84 (+364%, cache-thrashing)

Table 1: Inner-loop micro-benchmarks at $N=512$, 1 thread, 64 iters. `ji-order` is a textbook *wrong* answer: stride- N along the inner loop blows the L1 cache on every column. `restrict` is essentially free (gcc 11 already infers no-aliasing here. I checked the assembly).

The `ji` regression is the canonical Bryant-Hallaron- style memory-pattern result: row-major stored data wants row-major access. We’ll keep `ij-order` forever. And `restrict` buys us essentially nothing here. The stencil writes to `dst` and reads from `src`, which are different allocations, so the compiler had no aliasing to worry about in the first place. More on why this gets interesting once we add temporal-blocking scratch buffers.

2.2 The full sweep

We tried three full-sweep algorithms, in increasing order of complexity:

1. **Plain double-loop** (the **baseline** in every binary). Two nested loops, one big in-place pass, with the Jacobi ping-pong trick to make it parallelisable.
2. **Spatially blocked**. Tile the domain into $B \times B$ chunks (Figure 2), each thread owns a tile, scratch is just the tile + 1 cell of halo. Helps when the working set exceeds L2 (we’ll see this in Figure 4).
3. **Temporally blocked**. The shrinking trapezoid (Figure 2). Per-thread scratch is $2(B+2T)^2$ floats (two ping-pong buffers). We do T in-scratch sweeps before any DRAM round-trip.

For (3), the work multiplier is $((B+2T)/B)^2$. At $B=128$, $T=8$ that’s $(144/128)^2 \approx 1.27 \times$. 27% extra cell-updates compared to baseline. The bandwidth savings have to outpace this overhead for temporal blocking to win at all.

2.3 Loose ends

A few details that aren’t the optimisation priority but *do* show up in the timing:

- **$O(N)$ boundary copies**. We pass the boundary rows and columns through unchanged from `src` to `dst` every sweep. At $N=2048$ that’s $4N \cdot 4B = 32\text{ KB}$ of writes. Negligible vs. the $2N^2 \cdot 4B = 16\text{ MB}$ of interior data.
- **Per-tile scratch malloc inside the temporal kernel**. The OpenMP version does `malloc(2*(B+2T)*(B+2T)*sizeof(float))` once per thread per super-step. At small N this is real overhead. It shows up directly in the size-sweep result (Figure 4) where temporal *loses* to baseline at $N \leq 512$.
- **Convergence-norm reduction**. In `sor2d_rb.c` we accumulate $\|u^{(k+1)} - u^{(k)}\|_\infty$ in the inner loop; that’s an extra branch and an extra reduction per cell. It’s why the omega-sweep timings (2.5.2) are slower per-iter than the ping-pong tier.
- **Final result memcpy**. At the end of an even number of super-steps, the answer ends up in the wrong ping-pong buffer; we either swap or memcpy. We picked the cheap one (a pointer swap).

2.4 Stitching it together (real code)

Below is the actual temporal-blocking inner loop from `sor2d_omp.c`, with comments mapping each line back to 1.3’s pseudocode. Stripped of declarations for space; full source in the repository.

```
/* === one super-step: advance src -> dst by T Jacobi sweeps === */
/* === inside one #pragma omp parallel region, each thread takes === */
/* === one (B+2T)x(B+2T) scratch (sa,sb), iterates T times. === */

#pragma omp for collapse(2) schedule(static)
for (ti = 0; ti < ntiles_i; ti++) {
    for (tj = 0; tj < ntiles_j; tj++) {
        int i0 = 1 + ti*B, j0 = 1 + tj*B; /* tile origin in dst */
        int gi0 = i0 - T, gj0 = j0 - T; /* scratch origin */

        /* (1) load (B+2T)x(B+2T) scratch from src with clamp-to-edge */
        for (si = 0; si < S; si++) /* S = B + 2*T */
            for (sj = 0; sj < S; sj++) {
                gi = clamp(gi0+si, 0, N-1);
                gj = clamp(gj0+sj, 0, N-1);
                sa[si*S+sj] = src[gi*N + gj];
            }
        memcpy(sb, sa, sizeof(float)*S*S); /* both ping-pong buffers */

        /* (2) T sub-steps on the in-scratch trapezoid */
        A = sa; Bp = sb;
        for (t = 0; t < T; t++) {
            memcpy(Bp, A, sizeof(float)*S*S); /* carry boundary forward */
            for (si = 1+t; si < S-1-t; si++) /* shrinks by 1 each step */
                for (sj = 1+t; sj < S-1-t; sj++) {
                    float s = A[si*S+sj];
                    float nb = 0.25f*(A[(si-1)*S+sj] + A[(si+1)*S+sj]
                                     + A[si*S+sj-1] + A[si*S+sj+1]);
                    Bp[si*S+sj] = s - OMEGA*(s - nb); /* eq. (3) damped Jacobi */
                }
            tmp = A; A = Bp; Bp = tmp; /* in-scratch ping-pong */
        }

        /* (3) write central B x B back to dst */
        for (si = T; si < T+B; si++)
            memcpy(dst + (gi0+si)*N + (gj0+T),
                   A + si*S + T,
                   B*sizeof(float));
    }
}
```

The `collapse(2)` schedules tiles across threads as one flat ($\text{ntiles}_i \cdot \text{ntiles}_j$) work-list, so OMP can load-balance even when one of the dimensions is small. The `memcpy(Bp, A, ...)` in the inner t -loop is the “carry-the-frame-through” trick. It copies the previous sub-step’s whole scratch into the destination buffer so cells *outside* the shrinking trapezoid (the perimeter of stale halo, plus any cells that map to a global-boundary coordinate) are preserved without per-cell branches. The compute loop is then branch-free and `nvcc/gcc` both vectorise it cleanly.

2.5 Code and Results

2.5.1 Discussion of CPE

The size-sweep is the first place the cache-hierarchy story becomes visible (Figure 4, E02 in the bench).




Figure 4: 2D OMP size sweep at 8 threads, $B=128$, $T=8$, iters= 64. Median of 3 reps; shaded band is the min/max envelope. The L3→DRAM crossover lands between $N=1024$ and $N=2048$: baseline CPE *rises* from 0.464 to 0.667 as the working set exceeds 38 MB (one socket’s L3). Temporal stays at ~ 0.35 across $N=1024-4098$. It’s L2-resident on its scratch. And wins $1.93\times$ at $N=2048$. Below L3 (the two leftmost points) temporal *loses*: per-tile scratch malloc + redundant halo work cancel the savings on a working set the baseline already had in cache.

The L2 budget per Sapphire-Rapids core is roughly 2 MB, and the per-thread $(B+2T)^2$ scratch at $B=128$, $T=8$ is $144^2 \cdot 4B = 81$ KB. Comfortably L1, let alone L2. So when we say “temporal blocking earns its keep when the baseline goes off-chip”, we mean it literally: the temporal kernel’s working set is hot in L1 even at $N=4098$. The baseline’s working set at $N=2048$ is $2 \cdot 2048^2 \cdot 4B = 16$ MB. Fits L3 of one socket (38.4 MB), barely, but every sweep streams the whole thing through the cores. At $N=4098$ that’s 64 MB, which spills L3 entirely; the baseline runs at DRAM bandwidth and the temporal kernel runs at L1/L2 bandwidth. The $2\times$ ratio we measure is exactly the bandwidth ratio, give or take a constant for halo overhead.

Block size sweep. At fixed $N=2048$ we sweep B to find the temporal kernel’s spatial sweet spot (Figure 5, E03).

The cliff at $B \geq 384$ landed exactly where back-of- envelope predicts. We pick $B=128$ as the canonical choice for the rest of the experiments: it sits in the middle of the plateau, the scratch fits L2 with comfortable headroom, and $2046/128 = 16$ tiles per axis aligns cleanly so we don’t pay edge-tile irregularity.

Halo / T sweep. At fixed $B=128$ we sweep T to find the *temporal* sweet spot (Figure 6, E04).

The full table:

Strong scaling and decomposition. Two more figures worth quoting. Figure 7 shows the strong- scaling curves: temporal hits 91% of perfect scaling at 8 threads ($7.31\times$), baseline only 56% ($4.49\times$). At one thread temporal barely beats baseline (CPE 2.756 vs 2.792. $1.08\times$), because the Xeon’s hardware prefetchers + 38 MB L3 keep the single-thread baseline near




Figure 5: Spatial block-size sweep at $N=2048$, $T=4$, 8 threads. Baseline is flat (grey). B has no effect on the baseline kernel, which never sees a tile. Temporal has a wide plateau across $B \in [24, 256]$ at CPE ~ 0.39 , then *falls off a cliff* at $B \geq 384$. Red dashed line is the L2-fit threshold: $2(B+8)^2 \cdot 4B \leq 1.5 \text{ MB} \Rightarrow B \leq 297$. Beyond that, the per-thread scratch doesn't fit L2 and we pay L3 misses on our *own* scratch. Defeating the entire purpose.




Figure 6: Temporal depth T sweep at $N=2048$, $B=128$, 8 threads, iters= 96. Left axis: CPE (lower = faster). Right axis (red dashed): redundant halo work as a percentage, $((B+2T)^2/B^2 - 1) \cdot 100$. The two axes encode the project's central tradeoff: DRAM savings grow linearly in T , redundant halo work grows quadratically, and the unimodal product peaks at $T=8$ with CPE 0.332. The best CPU CPE we measured anywhere. $T=1$ loses to per-tile-malloc + memcpy overhead with no DRAM reuse; $T=32$ has 125% halo overhead and gives *everything* back.

streaming bandwidth and there's nothing for temporal blocking to recover. This is the *temporal blocking earns its keep only once parallelism stresses DRAM* story.


Figure 7 (E05) is the Lab-5-Part-4 redo: six lines, three decomposition modes crossed with two scheduling strategies, all at $N=2048$. Strip-persistent wins (CPE 0.255 at 8 threads, $7.89\times$ scaling); interleaved loses (false sharing on adjacent rows); 2D-block sits in the middle (no false sharing, but $2\times$ the halo of strip). The **spawn-vs-persistent** gap is largest for strip because strip's per-sweep work is smallest. $1.54\times$ slowdown for strip, $1.03\times$ for interleaved. Quantifying the per-sweep `pthread_create` cost is the next experiment.

Pthread spawn-overhead. Figure 8 (E07) is a direct Lab-6-Part-1b answer.

The slope difference is the per-sweep `pthread_create` cost: roughly $153\mu\text{s}$ at 4 threads, or

T	baseline CPE	temporal CPE	temp/base	halo overhead
1	0.504	0.739	$1.47\times$	3%
2	0.469	0.487	$1.04\times$	6%
4	0.464	0.393	$0.85\times$	13%
6	0.448	0.360	$0.80\times$	20%
8	0.444	0.332	$0.75\times$	27%
12	0.469	0.348	$0.74\times$	41%
16	0.412	0.378	$0.92\times$	56%
24	0.418	0.393	$0.94\times$	89%
32	0.401	0.461	$1.15\times$	125%

Table 2: Halo / T sweep at $N=2048$, $B=128$, 8 threads, iters= 96. Best temp/base ratio at $T=8$ ($0.75\times$. Temporal is 25% faster than baseline at the same iter count). At $T=1$ temporal is *worse*. It pays the per-tile-malloc + memcpy overhead with no DRAM-reuse benefit. At $T=32$ the halo overhead swallows the win. This is the central plot for the talk.



(a) Strong-scaling (E01). Temporal is closer to ideal.

(b) Pthreads decomposition $\times$ schedule (E05).

Figure 7: Left: strong scaling at $N=4098$, $B=128$, $T=8$. Right: 6 lines = 3 modes (strip / interleaved / 2D-block) $\times$ 2 schedules (persistent / spawn). Strip-persistent at 8 threads: CPE 0.255, $7.89\times$ scaling. Within 1.4% of perfect $8\times$. Interleaved bottoms out at $2.67\times$; the cause is false sharing on adjacent rows, since two threads ping-pong cache lines they share.

about $38\mu\text{s}$ per `pthread_create+join` pair. This is why OpenMP under the hood uses a persistent thread pool. It's not magic, it's just spawning once per program rather than once per parallel region.

The OMP-vs-pthreads surprise. At $N=2050$, iters=96, 8 threads, OMP-8t baseline measured CPE 0.667 (E11 also saw a similar config). Pthreads-strip-persistent at the same N measured CPE ~ 0.4 . *Pthreads is faster than OpenMP at the baseline sweep* on this machine. The plausible suspect is OpenMP's default `schedule(static)` chunking interacting with first-touch NUMA allocation. The buffers were allocated by thread 0, so on the second NUMA node every thread is reading remote-NUMA memory. We didn't chase this further (would have needed `numactl` + `libnuma` surgery), but it's the candidate root cause. An honest one to call out. I'd expected OMP to dominate everywhere.




Figure 8: Wall-time vs iters at $N=512$, 4 threads, mode=strip. Linear-fit $\text{time} = a + b \cdot \text{iters}$ for both schedules. Persistent slope: $81 \mu\text{s}/\text{sweep}$. The actual compute work. Spawn slope: $235 \mu\text{s}/\text{sweep}$. The slope difference of $153 \mu\text{s}$ is the per-sweep `pthread_create+join` overhead. **It's larger than the compute work itself.** Spawning fresh threads on every sweep is a $1.89\times$ tax on the kernel.

2.5.2 Discussion of Iteration Number and Error

The previous section was per-iter throughput. This one is iters- to-tolerance, which is where the whole *algorithmic* story lives. The classical Lab-5 figure: sweep ω at fixed N and plot iters-to-converge (Figure 9, E08).




Figure 9: Red-black GS-SOR omega U-curve at $N=128$, $\text{TOL}=10^{-5}$. 150 ω values in $[0.50, 1.99]$ in 0.01 steps. Theory $\omega_{\text{opt}} = 2/(1 + \sin(\pi/(N-1))) = \mathbf{1.9517}$ (red dashed) matches empirical minimum at $\omega = \mathbf{1.95}$ (green dotted) to **0.1%**. 254 iters at ω_{opt} vs 3,434 at $\omega=1$ (Gauss-Seidel). A **$13.5\times$ fewer-iteration win** from algorithmic tuning alone.

The cliff at $\omega > \omega_{\text{opt}}$ is sharp: $\omega=1.99$ already gives back most of the win (1,230 iters vs 254 at the optimum). Above $\omega=2$ the iteration is unconditionally divergent (the binary caps at `MAX_ITERS` and prints `DIVERGED`).

This is the project's biggest single algorithmic optimisation, and the bittersweet thing is the temporal-blocking kernels can't directly exploit it. Damped Jacobi at $\omega = 0.9$ lives on the *under-relaxation* side of the curve; $\omega > 1$ in the ping-pong form has spectral radius > 1 on large

grids and diverges unconditionally (we got bitten by this in the first session of the project, more on it in 4). The natural follow-up is a temporal-blocking variant of red-black GS-SOR. The shrinking trapezoid would have to be re-derived for paired half-sweeps. And that’s the most concrete future-work item below.

Convergence with the long-sweep FP32 question. At $N=128$, 254 iters at ω_{opt} and $\text{TOL}=10^{-5}$. The FP32 unit roundoff is $\sim 6 \cdot 10^{-8}$, and a worst-case forward-error bound after K stencil ops per cell is roughly $\sqrt{K} \cdot \varepsilon \sim 10^{-7} \cdot \sqrt{254 \cdot 5} \approx 10^{-6}$. Comfortably below TOL. At larger N or tighter TOL this bound starts to bite, and the responsible fix is either a Kahan-style compensated sum or moving to FP64. We didn’t need it.

2.6 Overall comments

1. **Temporal blocking only earns its keep in the bandwidth-bound regime.** Single-thread CPU was a wash ($1.08\times$); 8-thread + DRAM-spilling working set was $2\times$. This isn’t a flaw. It’s the textbook prediction.
2. **The T -sweet-spot is sharp.** At $T=8$ we get 25% faster than baseline; at $T=1$ we lose 47% to scratch-allocator overhead; at $T=32$ we lose 15% to halo work. Tune carefully.
3. **Strip wins for 2D pthreads.** Lab 5 already said so; we got $7.89\times$ scaling at 8 threads.
4. **OMP can be slower than pthreads at the baseline sweep on a NUMA box.** Suspect: first-touch + cross-socket threads. Honest negative. Diagnose, don’t paper over.
5. **ω_{opt} is real and matters more than any per-iter optimisation we tried on CPU.** $13.5\times$ fewer iters at $N=128$ from a single relaxation-parameter sweep. Combining this with temporal blocking is the #1 future-work item.

3 To the GPU!

The CPU side is done; we have a baseline running near DRAM bandwidth and a temporal kernel running near L2 bandwidth. The GPU has roughly **30** $\times$ more memory bandwidth (864 GB/s on the L40S vs ~ 280 GB/s per Xeon socket) and *way* more parallelism. The interesting question is whether temporal blocking’s relative advantage *shrinks* on the GPU, because a kernel that’s already near memory-bandwidth on CPU might be much closer to compute-bound on the GPU’s higher bandwidth. Spoiler: yes, it shrinks.

3.1 Memory allocation and transfer

The buffers we need on-device, at $N=2048$, FP32:

buffer	size	at $N=2048$	H→D once?
d_src	$N^2 \cdot 4\text{ B}$	16 MB	yes (init)
d_dst	$N^2 \cdot 4\text{ B}$	16 MB	no (alloc only)
shared-mem scratch (per block)	$2 \cdot \text{TILE}^2 \cdot 4\text{ B}$	8 KB	no (per kernel)
constants (N , ω , T)	negligible		yes (cudaMemcpyToSymbol)

Table 3: GPU memory layout at $N=2048$, FP32. The dominant transfer cost is the initial 16 MB host-to-device copy of **d_src**. At PCIe Gen4’s ~ 25 GB/s that’s ~ 0.6 ms or $\sim 2 \times 10^6$ cycles at 3.5 GHz. About the same wall time as a single GPU *baseline* kernel launch over 2048×2048 (~ 2.7 ms). After that, the iteration loop is fully on-device.

The convergence-norm reduction we’d need for a TOL-based stop isn’t in the temporal-blocking kernel because that kernel does T sweeps in shared memory before any global write. There’s no clean way to read out a partial residual mid-super-step. We just run a fixed iters and trust the math. This was a deliberate simplification; a real solver would need either a kernel-per-sweep form or a separate reduction-over-tiles kernel after each super-step. “It’s not the most significant saving”. But it does mean we can’t do adaptive convergence on the GPU in this implementation.

3.2 The parallelised hot loop on the GPU

Vocabulary watch. The CPU thread $\rightarrow$ CUDA mapping is non-obvious and caught me out the first time I tried it. CPU “thread” maps to CUDA *block* (each block has private shared memory and a local thread-team that can synchronise via `__syncthreads()`). CPU “block of work” (the $B \times B$ tile) maps to CUDA *tile* which we’ll call `TILE` throughout this section. The within-block warps are the equivalent of SIMD lanes inside a CPU thread. This vocabulary mismatch is worth telling the reader. I’ve seen otherwise solid people confused by it.

The temporal kernel’s structure is, in pseudocode:

```
__global__ void sor_temporal(const float *d_old, float *d_new,
                           int N, float omega, int tile, int halo)
{
    extern __shared__ float sshared[];
    float *sa = sshared;
    float *sb = sshared + tile*tile; /* two ping-pong buffers */

    int lx = threadIdx.x, ly = threadIdx.y;
    int gi = blockIdx.y * (tile - 2*halo) + 1 + (ly - halo);
    int gj = blockIdx.x * (tile - 2*halo) + 1 + (lx - halo);

    /* (1) clamp-to-edge load into both buffers */
    float v = d_old[clamp(gi)*N + clamp(gj)];
    sa[ly*tile + lx] = v; sb[ly*tile + lx] = v;
    __syncthreads();

    /* (2) HALO sub-steps; valid region shrinks 1 each side per step */
    int cur = 0;
    for (int t = 0; t < halo; t++) {
        bool in_trap = (lx >= 1+t && lx < tile-1-t
                       && ly >= 1+t && ly < tile-1-t);
        /* compute or carry-through */
        float out = in_trap ? compute_5pt(sa,sb,cur) : current_value(sa,sb,cur);
        __syncthreads();
        write_to_other_buffer(out, sa, sb, cur);
        __syncthreads();
        cur ^= 1;
    }

    /* (3) write central INTER x INTER cells back to d_new */
    if (lx >= halo && lx < tile-halo && ly >= halo && ly < tile-halo)
        d_new[gi*N + gj] = (cur == 0) ? sa[ly*tile + lx]
                                     : sb[ly*tile + lx];
}
```

The kernel uses dynamic shared memory ($2\text{TILE}^2 \cdot 4\text{B} = 8\text{KB}$ at `TILE=32`) so we can vary `TILE/HALO` at runtime without recompiling. At sm_89 the per-block shared-mem limit is 100 KB; we’re well under. The ridge is occupancy: at `TILE=32` each block has 1024 threads (32×32

the CUDA cap), which means only one block per SM. The warp scheduler has 32 warps to pick from, which is fine for hiding short-latency shared-mem dependencies but tight if the kernel ever stalls on a DRAM read. At `TILE=24` we get 576 threads per block = 18 warps, and the SM can fit two blocks simultaneously. We’ll see this play out empirically below.

Tuning sweep. Figure 10 (E09) is the (`TILE`, `HALO`) heatmap.




Figure 10: GPU 2D temporal kernel (`TILE`, `HALO`) sweep at $N=2048$, `iters=96`. 21 valid combos $\times$ 3 reps. Best at (`TILE=24`, `HALO=4`) with CPE 0.0088 – **$1.56\times$ over the per-launch baseline kernel** (CPE 0.0137). `TILE=8` is dead. Only 64 threads/block (2 warps) isn’t enough to hide DRAM latency. `HALO` too big collapses too: the empirical pattern is **`HALO` $\leq$ `INTER`/4**.

The same shrinking-trapezoid tradeoff as on CPU (Figure 6), at GPU scale. `HALO=4` wins consistently for `TILE` ≥ 16 ; the empirical rule is **`HALO` $\leq$ `INTER`/4**, where `INTER` is the central-output side. Best is (24, 4) with `INTER=16`, ratio 4:1.

3.3 Partitioning the rest of the algorithm

We had a real design choice between two GPU kernel structures:

1. **One kernel per stencil sub-step.** A thin `sor_sweep` kernel that does exactly one Jacobi sweep, host loop relaunches it `iters` times. Clean and easy to add a residual reduction between launches; matches the Lab-7 form exactly. But it pays $N^2 \cdot \text{iters}$ DRAM round-trips and one $\sim 30 \mu\text{s}$ launch each time.
2. **One temporal-blocked kernel** that holds a $(\text{TILE})^2$ shared-memory tile and does `HALO` sub-steps inside one launch. Fewer DRAM trips, fewer launches, but mid-super-step residual checks are awkward and the kernel is bound by shared-mem capacity.

We implemented both (`sor_sweep` for the baseline, `sor_temporal` for the temporal kernel) and run them side-by-side in the same binary. Letting the host swap pointers between launches is also *ping-pong* (CUDA-side this time), so both kernels share the same overall structure. The temporal kernel just amortises 4 sweeps across one launch instead of one. Shared-memory budgeting at `TILE=24`, `HALO=4`: $2 \cdot 24^2 \cdot 4 \text{ B} = 4.5 \text{ KB}$ per block, well under the 100 KB `sm_89` cap.

3.4 Running it end-to-end

Figure 11 is the GPU N-sweep (E10).




Figure 11: 2D GPU baseline vs temporal across N , iters= 96, defaults TILE=32, HALO=4. Median of 3 reps; shaded band is min/max envelope. At small N the baseline is launch-overhead-dominated ($96 \text{ launches} \times \sim 30 \mu\text{s} \approx 3 \text{ ms}$ is most of its wall time at $N=258$); the $17\times$ “speedup” there is mostly launch-count reduction ($96 \rightarrow 24$), not bandwidth. At $N \geq 1024$ the steady-state speedup is a more honest **$1.5\text{--}3\times$** . At $N=4098$ both kernels degrade because two FP32 buffers (128 MB) exceed the L40S’s $\sim 96 \text{ MB}$ L2. The GPU equivalent of the CPU L3 cliff in Figure 4, two orders of magnitude larger. Best 2D GPU CPE = **0.0089** at $N=3074$.

N	baseline ms	temporal ms	baseline CPE	temporal CPE
258	2.7	0.16	0.395	0.023
514	1.8	0.27	0.096	0.014
1026	2.4	0.83	0.030	0.010
1538	3.7	1.86	0.018	0.0091
2050	5.7	3.83	0.014	0.0090
3074	9.5	8.18	0.010	0.0089
4098	43.6	18.4	0.026	0.011

Table 4: 2D GPU N-sweep, iters= 96. CPE values reported at CPNS= 2.0 (lab convention), so the absolute scale is approximate on this 2.5–4.1 GHz Ada part, but the *ratios* between runs are exact.

3.5 Discussion of GPU results

Oi! Well, the GPU temporal/baseline ratio isn’t great. On CPU we got $1.93\times$ (E02 at the sweet spot $N=2048$); on GPU we get only $1.56\times$ (E09 best config) or even $1.16\times$ at $N=3074$ where both kernels are near peak. Why so much smaller?

The L40S has $\sim 864 \text{ GB/s}$ of HBM-class memory bandwidth, an $\sim 96 \text{ MB}$ L2, and $\sim 100 \text{ KB}$ shared memory per SM. At $0.75 \text{ FLOP/byte AI}$, the bandwidth-bound ceiling is $0.75 \cdot 864 = 648 \text{ GFLOP/s}$. We measured GPU baseline at $\sim 1.36 \text{ TFLOP/s}$ (CPE $0.0103 \rightarrow$ throughput from (5)). That’s *twice* the DRAM-bound roof, which means the baseline kernel is already running mostly out of L2 ($\sim 4.5 \text{ TB/s}$ sustained), not DRAM. When the baseline is already cache-resident, temporal blocking has nothing to recover – the L2 is a 96 MB cache, our $2N^2$ buffer at $N=2050$ is 33 MB , the working set fits. We only see the genuine memory-saving win at $N=4098$ (128 MB working set, exceeds L2) where the ratio re-opens to $2.36\times$.

Apples to apples. Quoted speedup against the *serial-CPU baseline* is $310\times$ (CPE $2.792 \rightarrow 0.009$). That’s the headline number people put on slides, but it’s not the honest one. Apples-to-apples comparisons:

- GPU-temporal vs OMP-temporal (the two best from each tier): $0.332/0.009 \approx 37\times$. This is the legitimate “GPU is better than 8-core CPU” number for this workload at this N .
- GPU-temporal vs GPU-baseline (the optimization within the tier): $1.56\times$. This is the legitimate “shared-memory temporal blocking helps” number.
- GPU-baseline vs serial-CPU-baseline (just the architecture): $186\times$. This is the legitimate “the L40S is much better than one Xeon thread” number.

The $310\times$ is just $186 \cdot 1.56 \approx 290$ (rounded). It’s the product of two effects, presented on its own it overstates the within-architecture optimization. We’ll quote it in the headline but be careful in the accompanying text.

The 3D GPU honest negative. We also implemented a 3D temporal kernel (`sor3d_gpu.cu`, TILE=8, HALO=2) and ran it alongside the 3D baseline. Result in Figure 12: the 3D temporal *loses* to the 3D baseline at $N \geq 194$. The arithmetic is unforgiving: with TILE=8, HALO=2, INTER=4, the halo- to-interior ratio in 3D is $(\text{TILE}/\text{INTER})^3 - 1 = (8/4)^3 - 1 = 7\times$ redundant work. An order of magnitude worse than the $0.78\times$ overhead the 2D kernel pays at the same TILE/HALO ratio. The 2.5D-streaming variant (Mickevicius 2009 – slide a 3-z-plane window through shared memory instead of loading a cubic tile) is the canonical fix and is on the future-work list.




Figure 12: 3D GPU baseline (grey) vs temporal (red) across N . At $N \geq 194$ the temporal kernel loses. The 3D halo overhead ($7\times$ redundant work at TILE=8/HALO=2) eats the DRAM savings. This is the project’s headline negative finding; the fix is well- known (2.5D streaming) but wasn’t implemented.

Cross-tier headline (E11). Putting it all on one axis at $N=2050$, iters= 96, 8 threads: At CPNS= 2.0, GPU temporal’s 0.009 CPE corresponds to ~ 4.5 ps per element on the L40S. About one floating-point op per nanosecond per cell, which puts us near the machine-balance line for a memory-bound 5-point stencil at FP32. There’s no easy $2\times$ left at this point without changing the algorithm. Multigrid or 2.5D-streaming, both noted as future work.

tier / kind	CPE	vs serial baseline
serial-CPU baseline	2.792	1.00×
serial-CPU temporal	2.586	1.08×
pthread-8t-strip persistent	0.399	7.0×
OMP-8t baseline	0.667	4.2×
OMP-8t temporal	0.332	8.7×
GPU baseline	0.015	186×
GPU temporal	0.009	310×

Table 5: Cross-tier CPE summary (E11). “× vs serial baseline” is computed against **serial-CPU baseline**; the inside-tier comparisons (GPU temporal vs GPU baseline = 1.56×, GPU temporal vs OMP temporal = 37×) are the honest ones.

4 Closing Thoughts

What worked. Temporal blocking on the bandwidth-bound CPU regime was textbook: 1.93× at $N=2048$ where the working set first spills out of L3. The shrinking-trapezoid scheme in Figure 2 was correct by construction (every CPU config was bit-identical to its baseline) and the unimodal T -curve in Figure 6 traced exactly the shape predicted by the linear-savings-vs-quadratic-overhead analysis. GPU shared-memory tiling was also clean: 1.56× over the per-launch baseline at the right (TILE, HALO) sweet spot, and the kernel ran near 1.5 TFLOP/s on the L40S for a 0.75-AI stencil.

What didn’t. Three things in increasing severity:

1. **Single-thread CPU temporal was a wash (1.08×).** The per-tile-malloc + redundant halo work cancelled the bandwidth savings on a single core where DRAM was already lightly loaded. Not a flaw. Temporal blocking earns its keep when the baseline is bandwidth-bound, and the single-thread baseline isn’t. But worth being upfront about.
2. **OMP baseline (CPE 0.667) was worse than pthreads baseline (CPE 0.399)** at $N=2048$, 8 threads. Best guess is first-touch NUMA + the OMP runtime spreading threads across both sockets; the pthreads version manually pinned threads with `OMP_PROC_BIND=close` `OMP_PLACES=cores` and it shouldn’t have mattered, but it apparently did. I didn’t chase this with `numactl`. If I could beg a guess, it could be that, or it could be the OMP scheduler introducing slight stagger between threads that the barrier-only pthreads loop doesn’t. Honest unanswered.
3. **3D GPU temporal lost to baseline at $N \geq 194$.** Cubic tile $\Rightarrow$ surface-to-volume ratio that gets worse cubically; with TILE=8/HALO=2 the work multiplier was 7× and there’s no DRAM saving that can recover that. The fix (2.5D streaming) was on the to-do list and didn’t make the cut.

Precision and numerical issues. The biggest ones we hit, in chronological order: (1) Lab-7’s $\omega=1.85$ in damped- Jacobi-ping-pong form is unconditionally divergent on large grids (spectral radius > 1). Early runs produced `inf/NaN` that the `max_abs_diff` validator silently masked because `NaN > 0` is false. Fixed by switching to $\omega=0.9$ in the temporal kernels and reserving $\omega>1$ for the red-black GS variant; the `max_abs_diff` now flags non-finite values explicitly. (2) FP32 round-off accumulates as $\sqrt{N} \cdot \text{iters} \cdot \varepsilon_{\text{FP32}} \approx 10^{-5}$ at the configs we ran; comfortably below TOL but would become a problem at larger N or tighter TOL. (3) Red-black ordering

shifts the iteration matrix’s spectral radius slightly relative to scan-order GS, which makes ω_{opt} a function of grid size (formula (4)); we verified the N -dependence in 2.5.2 matches the formula.

Future directions. Three concrete next steps, in priority order:

1. **Red-black GS-SOR temporal kernel.** Combine the $13.5\times$ algorithmic win from Figure 9 with the $1.56\times$ per-iter win from Figure 10. Re-derive the shrinking trapezoid for paired half-sweeps; probably a couple of days of careful indexing. Expected total speedup: $20\times$ on the convergence side $\times 1.5\times$ on the per-iter side $\approx 30\times$ over the current GPU temporal.
2. **2.5D-streaming for the 3D GPU temporal kernel** (Micikevicius 2009). Slide a 3-z-plane window through shared memory instead of loading a cubic tile; turns the 3D’s $7\times$ halo overhead into 2D’s $0.78\times$. Should flip Figure 12’s temporal/baseline ratio to > 1 everywhere.
3. **Multigrid V-cycle.** The proper way to amplify SOR’s per-iter speedup with a $5\text{--}20\times$ fewer-iterations win on top. V-cycle restriction/prolongation is straightforward; each level’s smoother is just our existing red-black SOR. This is also the right sanity check on “iterations to tolerance”. Multigrid solvers converge in a small constant number of V-cycles regardless of N , so any residual-vs-iterations curve we plot for SOR can be compared against multigrid as the true asymptotic baseline.
4. *(Easter egg) Mixed precision.* Keep an FP64 residual accumulator in the outer iteration loop and FP16/ BF16 inside the inner stencil; iterative-refinement style. ConvStencil (PPoPP 2024) showed $2\text{--}5\times$ wins on Tensor Cores for related stencil shapes. We’re nowhere near Tensor-Core-ready; commented out the framework but left it in the source as a fun easter egg.

5 Compilation instructions

```
# CPU binaries (all *.c sources, including the OMP and pthreads variants)
gcc -O1 -std=gnu11 -fopenmp src/sor2d_omp.c -o build/sor2d_omp -lrt -lm
gcc -O1 -std=gnu11 -pthread src/sor2d_pth_decomp.c -o build/sor2d_pth_decomp \
    -lpthread -lrt -lm

# GPU binaries (need cuda module loaded)
module load cuda/12.8
nvcc -O2 src/sor2d_gpu.cu -o build/sor2d_gpu
nvcc -O2 src/sor3d_gpu.cu -o build/sor3d_gpu

# Or just 'make && make gpu'. The project-root Makefile does the same.
```

Why `-O1` and not `-O3`: matches the EC527 lab convention (verified by `grep -r "gcc -O" EC527/Lab*`). The lab CPE numbers in Bryant–O’Hallaron all assume `-O1` to keep the compiler from auto-vectorising and complicating the analysis. We’d absolutely use `-O3 -march=native -ffast-math` in production. A quick check showed `gcc-11` auto-vectorises the inner stencil into `AVX-512 vfmadd231ps` at `-O3` and roughly halves the CPU-baseline CPE. That’s a free $2\times$ we chose not to take, for course-style consistency.

Preprocessor knobs the reader can flip in the source:

- `N`. Grid side (set at runtime via `argv[1]`).
- `B`, `T`. Spatial / temporal block sizes (`argv[3]`, `argv[4]` for the OMP/CPU binaries).
- `OMEGA`. Relaxation factor. Default 0.9 (damped Jacobi); set to $\omega_{\text{opt}}(N)$ for `sor2d_rb`.
- `TOL`. Convergence tolerance. 10^{-5} default.
- `ITERS`. Fixed iter count for the temporal binaries (must satisfy `ITERS % T == 0`).

- GPU runtime: `--tile T` and `--halo H` on `sor2d_gpu` and `sor3d_gpu`.

List of files

Source files in `src/`:

- `sor2d_cpu.c`. Single-thread 2D baseline + temporal kernel.
- `sor3d_cpu.c`. Single-thread 3D, 7-point stencil.
- `sor2d_omp.c`. OMP 2D baseline + temporal.
- `sor3d_omp.c`. OMP 3D.
- `sor2d_pth.c`. Pthreads 2D, strip decomposition, barrier-per-sweep.
- `sor2d_pth_decomp.c`. Pthreads decomposition study (3 modes $\times$ 2 schedules) used in Figure 7.
- `sor2d_pth_temporal.c`. Pthreads 2D with temporal blocking (mirror of the OMP version).
- `sor3d_pth_temporal.c`. Pthreads 3D temporal.
- `sor3d_omp_part.c`. 3D OMP partitioning study (slab/pencil/cube).
- `sor2d_rb.c`. Red-black GS-SOR with ω sweep used in 2.5.2.
- `sor2d_gpu.cu`. 2D GPU baseline + shared-memory temporal kernel.
- `sor3d_gpu.cu`. 3D GPU.

Result data in `results/exp/E0{0--11}/`, one `notes.md` per experiment, raw run logs in `raw_*.txt`, derived tables in `data.csv`, plots in `report_images/` (this report references those PNGs only). `results/STATUS.md` is the one-glance dashboard.